

A Statistical Audit of LLM Calibration Heterogeneity Across Knowledge Domains

Jennifer Zhang

John Wright

STAT 215B, Spring 2026

Abstract

Large language model (LLM) calibration is nearly always reported as a single aggregate scalar, obscuring whether miscalibration is diffuse or concentrated in specific knowledge domains. We develop a statistically rigorous audit framework and apply it to MMLU [Hendrycks et al., 2021] using token-level log probabilities from seven instruction-tuned models spanning two families and six parameter scales: GPT-4o-mini, Llama-3-8B-Instruct, and Qwen-2.5-{0.5B, 1.5B, 3B, 7B, 14B}-Instruct. Our framework combines a three-level mixed-effects model, isotonic and kernel-smoothed reliability curves, empirical Bayes shrinkage, and Benjamini–Hochberg FDR correction. Across the seven models, between-subject calibration variance accounts for 1.7–7.9% of the total, and 55–57 of 57 subjects are flagged as significantly miscalibrated at FDR $\alpha = 0.05$. Subject-level ECE rankings are highly correlated for the six larger models (Spearman $\rho \in [0.80, 0.94]$, all 15 pairs); the smallest model (Qwen-2.5-0.5B) shows extreme positional bias (predicting D 59% of the time) that produces a qualitatively different miscalibration pattern uncorrelated with the others. The same subjects are robustly worst-calibrated across the larger models (college-level quantitative STEM, niche factual recall, structured ethical/legal reasoning) and the same subjects are robustly best-calibrated (high-school-level Social Sciences and applied/factual subjects). Within the Qwen family across four matched-precision scale points (1.5B/3B/7B/14B), the entropy coefficient grows monotonically with scale ($-0.004 \rightarrow +0.006 \rightarrow +0.029 \rightarrow +0.040$, last two $p < 0.001$), and between-subject ICC jumps from 1.5B to 3B ($0.017 \rightarrow 0.051$) and is essentially flat thereafter (0.049 at 7B, 0.059 at 14B). The mean-overconfidence intercept is non-monotonic, peaking at Qwen-3B (0.311) where strong confidence priors have been acquired ahead of capability gains, and declining at 7B and 14B as accuracy catches up. This is a marginal pattern consistent with prior aggregate findings on RLHF-induced confidence saturation [Kadavath et al., 2022, Nakkiran et al., 2025].

1 Introduction

Large language models are increasingly deployed in high-stakes settings—healthcare, law, education—where the reliability of model-expressed confidence is as important as accuracy itself. A well-calibrated model is one whose stated confidence tracks its empirical accuracy: if a model assigns 80% probability to an answer, it should be correct 80% of the time. Miscalibration, particularly overconfidence, means that users who rely on model confidence as a decision signal will be systematically misled.

A substantial literature has studied LLM calibration. Guo et al. [2017] established the modern empirical study of neural network calibration, popularizing expected calibration error (ECE) as a standard metric and demonstrating the effectiveness of temperature scaling as a post-hoc correction. Kadavath et al. [2022] showed that for multiple-choice tasks, token-level log probabilities provide a meaningful calibration signal, and that base (pre-RLHF) models exhibit better-calibrated

token probabilities than instruction-tuned counterparts. Xiong et al. [2024] provided a systematic benchmark of confidence elicitation methods, finding that all methods “struggle in challenging tasks, such as those requiring professional knowledge”—suggesting calibration quality may vary by domain, though the authors do not investigate this directly. Jiang et al. [2021] similarly found that question-level features such as answer entropy predict calibration quality, a pattern we examine formally here.

What is missing is a *statistical audit* of that heterogeneity. To our knowledge, the existing literature reports subject-level ECE without (i) stabilizing estimates against large differences in per-subject sample size, (ii) formally testing which subjects are robustly miscalibrated after correcting for multiple comparisons, (iii) decomposing how much total calibration variance is attributable to domain versus subject versus individual question, or (iv) characterizing subjects by their *pattern* of miscalibration rather than its scalar magnitude.

This framing is analogous to the gap between per-group accuracy reporting and a rigorous fairness audit [Barocas et al., 2023]: the fairness literature long ago recognized that single-number summaries are insufficient for subgroup auditing, and we import that logic into calibration evaluation. The downstream implication of the audit is structural: if miscalibration is diffuse, a single global recalibration (e.g., temperature scaling) is a reasonable fix; if it is concentrated in specific subject clusters and those clusters are stable across models, then subject-aware recalibration becomes a structurally appropriate scoping choice—a position that has been validated empirically by Luo et al. [2025], who fit subject-specific temperature parameters per MMLU subject. Our contribution is the diagnostic arm of that picture: stable audit findings establish whether the structural assumption behind a subject-aware correction is warranted in the first place.

Our central scientific questions, restated:

1. **Where** is calibration good and bad? Variance decomposition, FDR-corrected subject rankings, and a within-MMLU contrast between well-calibrated and poorly-calibrated subject clusters.
2. **How consistent** is the answer across models? Pairwise Spearman correlations of subject-level ECE across seven models.
3. **What mechanisms** produce the observed heterogeneity? Question-level predictors, plus a within-Qwen scaling sweep across four matched-precision checkpoints (1.5B, 3B, 7B, 14B).

2 Related Work

Calibration of neural networks. Guo et al. [2017] showed that modern neural networks are systematically overconfident and that temperature scaling largely corrects this at the aggregate level, with Zadrozny and Elkan [2002] earlier establishing isotonic regression as an alternative nonparametric recalibration method. Kadavath et al. [2022] extended calibration evaluation to LLMs, establishing that token-level probabilities on multiple-choice tasks carry a meaningful calibration signal. Nakkiran et al. [2025] and Yaldiz et al. [2026] study how instruction tuning affects calibration, finding that RLHF-style fine-tuning tends to degrade calibration relative to base models, an effect also documented in Ouyang et al. [2022].

Domain-level heterogeneity. Xiong et al. [2024] is the closest prior work in spirit, benchmarking confidence elicitation across task types. Jiang et al. [2021] showed that question-level features predict calibration on QA tasks. Luo et al. [2025] fits subject-specific temperature parameters for each of MMLU’s 57 subjects, implicitly acknowledging that a global correction is insufficient; Tan

et al. [2026] uses base model signals for unsupervised recalibration. Neither work performs a formal statistical audit of subject-level heterogeneity.

Multiple testing and empirical Bayes. The Benjamini–Hochberg procedure [Benjamini and Hochberg, 1995] controls the false discovery rate and is standard in high-dimensional inference. Empirical Bayes shrinkage of group-level estimates [e.g., Efron, 2012] is standard in multilevel modeling via best linear unbiased predictors (BLUPs). We use the multilevel modeling framework of Raudenbush and Bryk [2002] to decompose calibration variance across levels of the MMLU hierarchy.

3 Data

MMLU. We use the Massive Multitask Language Understanding benchmark [Hendrycks et al., 2021], accessed via the `cais/mmlu` dataset on HuggingFace. MMLU comprises 14,042 four-option multiple-choice questions across 57 subjects nested in four broad domains: STEM, Social Sciences, Humanities, and Other. Subject sizes range from 100 to 1,534 questions (median 163). The three-level hierarchy—questions $\rightarrow$ subjects $\rightarrow$ domains—is natural for the proposed analysis. We use the test split throughout.

Models. We query seven instruction-tuned models spanning two families and six parameter scales. **GPT-4o-mini** is queried via the OpenAI Batch API at temperature 0 (`max_tokens=1`, `top_logprobs=20`); the choice of 20 (rather than the API default of 5) ensures all four answer tokens are captured for virtually every question. **Llama-3-8B-Instruct** and **Qwen-2.5 {0.5B, 1.5B, 3B, 7B, 14B}-Instruct** are run locally on a Colab GPU. The 3B, 7B, 8B, and 14B models use 4-bit quantization (BitsAndBytes; Dettmers et al. 2022); the 0.5B model also uses 4-bit, and the 1.5B model is run twice (once at 4-bit, once at fp16) so the precision effect can be examined directly.¹ For all seven models we pre-fill the assistant turn with “The answer is” before extracting logprobs at the next-token position. Test-set accuracies are 0.749 (GPT-4o-mini), 0.606 (Llama), 0.405 (Qwen-0.5B), 0.571 (Qwen-1.5B fp16), 0.601 (Qwen-3B), 0.685 (Qwen-7B), 0.757 (Qwen-14B). Qwen-0.5B exhibits extreme positional bias (predicting D 59% of the time vs. A 5.6%), discussed where it affects results.

4 Methods

4.1 Confidence Elicitation and Outcome

For each question, the model’s confidence $c_i \in (0, 1]$ is the maximum of the softmax-renormalized distribution over the four answer tokens A/B/C/D, equal to $P(\hat{y} \mid \hat{y} \in \{A, B, C, D\})$ —the standard approach in calibration research [Kadavath et al., 2022]. The primary outcome is the *signed calibration gap*, $\text{gap}_i = c_i - y_i$ with $y_i \in \{0, 1\}$ indicating correctness; it is positive under overconfidence, negative under underconfidence, and supports regression-based inference that binning-based ECE cannot. Note that since $y_i \sim \text{Bernoulli}(p_j)$, residual variance is highest near 50% accuracy, so the homoskedasticity assumption of the mixed model below is imperfect.

¹The fp16 1.5B run was added after we observed that the 4-bit 1.5B run had unusual positional bias (predictions skewed toward B at 36%) and accuracy 0.054 below the published full-precision figure. The fp16 run brought both numbers in line with the published figures and is used as the canonical 1.5B in the main analysis. The 4-bit version is preserved for the precision-effect appendix.

Question-level covariates. We extract four features from each question, all standardized before fitting: **word count** of the question stem; **max option length** (word count of the longest answer choice); **negation indicator** (*not, except, least, never, no*); and **entropy** of the softmax distribution over A/B/C/D in nats, $H_i = -\sum_k p_{ik} \log p_{ik}$, ranging from 0 to $\log 4 \approx 1.386$.

4.2 Nonparametric Calibration Curves

For each subject we estimate a reliability curve using isotonic regression [Zadrozny and Elkan, 2002] rather than equal-width binning, with an unconstrained Gaussian kernel smoother ($\sigma = 1.5$ bins) overlaid as a monotonicity-assumption check. ECE per subject is the weighted mean absolute deviation between the reliability curve and the diagonal, using 20 equal-width confidence bins.

4.3 Multilevel Regression and Variance Decomposition

We model the question-level calibration gap across all 14,042 questions using a three-level mixed model with questions nested in subjects nested in domains:

$$\text{gap}_{ijk} = \mu + \mathbf{x}_{ijk}^\top \boldsymbol{\beta} + u_k + v_{jk} + \varepsilon_{ijk}, \quad (1)$$

where $\mathbf{x}_{ijk}$ are standardized question-level covariates, $u_k \sim \mathcal{N}(0, \sigma_{\text{domain}}^2)$ is a domain random intercept, $v_{jk} \sim \mathcal{N}(0, \sigma_{\text{subject}}^2)$ is a subject-within-domain random intercept, and $\varepsilon_{ijk} \sim \mathcal{N}(0, \sigma_{\varepsilon}^2)$ is residual noise. The model is estimated by restricted maximum likelihood (REML) using `statsmodels.MixedLM` [Raudenbush and Bryk, 2002].

The fixed effects $\boldsymbol{\beta}$ characterize which question features are associated with miscalibration across the full dataset, interpreted as descriptive associations rather than causal effects. The variance components ($\sigma_{\text{domain}}^2, \sigma_{\text{subject}}^2, \sigma_{\varepsilon}^2$) answer RQ1 via an intraclass correlation (ICC) decomposition: the fraction of total calibration heterogeneity attributable to each level of the hierarchy. Subject-level shrunk calibration estimates are the fitted random effects $\hat{v}_{jk}$ (BLUPs), which implicitly borrow strength across subjects in proportion to sample size.

4.4 Multiple Testing and Cross-Model Replication

For each of the 57 subjects we test H_0 : mean calibration gap = 0 via a two-sided one-sample t -test on the question-level gaps, controlling the false discovery rate at $\alpha = 0.05$ via Benjamini–Hochberg [Benjamini and Hochberg, 1995]. The full pipeline is run on each of the seven models, and subject-level ECE rankings are compared across all 21 model pairs with Spearman rank correlation and bootstrap CIs (2000 replicates, seed 42).

5 Results

5.1 Variance Decomposition (RQ1)

Table 1 reports the ICC decomposition from the three-level mixed model for the seven main models. Across all seven, the bulk of calibration-gap variance lies at the question level (92.1–98.3%), and domain-level variance collapses to zero. Between-subject ICCs span nearly a factor of five: GPT-4o-mini concentrates 7.9% of variance at the subject level, Qwen-14B 5.9%, Qwen-3B 5.1%, Qwen-7B 4.9%, and Llama and the smaller Qwen models cluster around 1.7–2.0%.

The within-Qwen scaling sweep across four matched-precision checkpoints (1.5B fp16, 3B 4-bit, 7B 4-bit, 14B 4-bit) shows ICC making a substantial jump from 1.5B to 3B (0.017 $\rightarrow$

Table 1: Variance decomposition of the calibration gap. Domain-level variance is estimated at zero for every model; $\text{ICC}_{\text{subject}}$ is the share of total variance attributable to subject-level heterogeneity.

Model	σ_{domain}^2	$\sigma_{\text{subject}}^2$	σ_{ε}^2	$\text{ICC}_{\text{subject}}$
GPT-4o-mini	0.000	0.013	0.150	0.079
Llama-3-8B-Instruct	0.000	0.004	0.190	0.020
Qwen-2.5-0.5B-Instruct	0.000	0.004	0.224	0.018
Qwen-2.5-1.5B-Instruct	0.000	0.003	0.197	0.017
Qwen-2.5-3B-Instruct	0.000	0.011	0.203	0.051
Qwen-2.5-7B-Instruct	0.000	0.009	0.177	0.049
Qwen-2.5-14B-Instruct	0.000	0.010	0.156	0.059

0.051) and remaining elevated at the larger scales (0.049 at 7B, 0.059 at 14B). The discontinuity around 1.5B→3B suggests that subject-level concentration of miscalibration emerges around 2–3B parameters, then increases gradually; combined with the parallel monotonic increase in the entropy coefficient (Section 5.3), this is the cleanest within-family scaling signal in our data.

The domain-level variance is estimated at zero for every model. With only four domain groups, the likelihood is flat with respect to the domain variance component, and the estimate collapses to the boundary. We therefore report total between-subject variance without a clean domain/subject partition, and treat domain as a descriptive grouping rather than a modeled random effect.

5.2 Calibration Curves by Subject

Figure 1 contrasts reliability curves for the five most and five least miscalibrated subjects (by ECE) for GPT-4o-mini. Curves are masked outside the observed 1st–99th percentile of confidence per subject to avoid displaying isotonic extrapolation: GPT-4o-mini’s confidence is so concentrated above 0.7 that any plotted curve below this region would be pure extrapolation. Within the observed support, the worst-calibrated subjects show the characteristic shape of an overconfident model—empirical accuracy lying below the diagonal, sometimes sharply—while the best-calibrated subjects track the diagonal closely. Both shapes are present in the same model on the same evaluation, illustrating the heterogeneity that aggregate ECE conceals.

5.3 Predictors of Miscalibration (RQ3)

Table 2 reports the fixed effects from the multilevel model across the seven main models. Intercepts confirm large mean overconfidence in every model (0.179–0.311). Within the Qwen family, the intercept is *non-monotonic* in scale: it rises sharply from 1.5B (0.179) to 3B (0.311), then declines back at 7B (0.233) and 14B (0.210). Mean accuracy is monotonic (0.571 → 0.601 → 0.685 → 0.757), so the inverted-U intercept reflects the confidence outpacing accuracy at intermediate scale, then accuracy catching up at larger scales.

Figure 2 shows the corresponding forest plot. Word count is a positive significant predictor for GPT-4o-mini and the larger Qwens (7B, 14B), inert for Llama and the small Qwens (0.5B, 1.5B, 3B). The most informative pattern is the entropy coefficient.

Entropy: a sign progression with scale. A naïve prediction is that high option-distribution entropy reduces the calibration gap, since top-1 confidence is mechanically lower when probability mass spreads across options. Yet the entropy coefficient varies systematically across models and




Figure 1: GPT-4o-mini reliability curves. **Top:** five most miscalibrated subjects by ECE. **Bottom:** five best-calibrated subjects. Same conventions in both: solid line is isotonic regression, plotted only on observed support (1st–99th percentile per subject); markers are 20-bin empirical accuracy sized by bin count; dashed line is the calibration diagonal.

Table 2: Fixed effects from the three-level mixed model for all seven models. Continuous covariates standardized. Significance: *** $p < 0.001$, ** $p < 0.01$, * $p < 0.05$, $^\dagger p < 0.10$.

Predictor	GPT	Llama	Q-0.5B	Q-1.5B	Q-3B	Q-7B	Q-14B
Intercept	0.196***	0.197***	0.209***	0.179***	0.311***	0.233***	0.210***
Word count	0.031***	-0.005	-0.013*	0.003	0.010	0.019**	0.023***
Max option length	0.001	-0.001	-0.028***	0.005	-0.008 †	-0.007	-0.007
Negation	-0.001	0.007	0.016	-0.010	-0.014	-0.001	0.005
Entropy	0.016***	0.008 †	-0.063***	-0.004	0.006	0.029***	0.040***

scales: strongly negative for Qwen-0.5B (-0.063 , $p < 0.001$), near-zero for Qwen-1.5B and Qwen-3B (-0.004 and $+0.006$, both n.s.), and positive and growing across larger models (Llama $+0.008$, GPT-4o-mini $+0.016$, Qwen-7B $+0.029$, Qwen-14B $+0.040$, all $p < 0.10$ or smaller). Within the Qwen family at matched precision (1.5B fp16, 3B, 7B, 14B), the coefficient grows monotonically across all four scale points: $-0.004 \rightarrow +0.006 \rightarrow +0.029 \rightarrow +0.040$. Figure 3 traces the mechanism. Within each model we partition questions into entropy quartiles and report mean confidence (solid) and mean accuracy (dashed); the shaded area is the calibration gap.

Within the Qwen family the pattern is one of confidence growing faster than accuracy as items become harder, rather than accuracy collapsing on hard items: at 1.5B confidence and accuracy move roughly together across quartiles (gap $+0.09$ at Q1, $+0.13$ at Q4), but at 7B and 14B confidence remains near saturation across the first three quartiles while accuracy still drops with entropy (Q1 vs. Q4 gaps: $+0.05$ vs. $+0.37$ at 7B; $+0.03$ vs. $+0.46$ at 14B). This is consistent with prior reports of RLHF-induced confidence saturation [Kadavath et al., 2022, Nakkiran et al., 2025, Yaldiz et al., 2026]. We treat the within-family comparison as observational—four checkpoints, one family—and read Qwen-0.5B’s strongly negative entropy coefficient as an artifact of its positional bias rather than a fourth scale point.

5.4 FDR-Corrected Subject Rankings

Figure 4 shows all 57 subjects sorted by mean calibration gap for each of the seven main models, with FDR-rejected subjects highlighted. At $\alpha = 0.05$: GPT-4o-mini, Llama, Qwen-3B,




Figure 2: Fixed effect estimates with 95% confidence intervals for all seven models. Estimates are on the standardized calibration gap scale. The entropy coefficient shows a striking sign progression across the Qwen scaling sweep.

Qwen-7B, and Qwen-14B reject all 57 subjects; Qwen-1.5B (fp16) rejects 56/57 (the lone exception is `jurisprudence`, $\bar{g}\bar{a}p = 0.051$, $p_{adj} = 0.18$); Qwen-0.5B rejects 55/57 (the exceptions are `high_school_statistics` and `formal_logic`).

The substantive finding is the *magnitude* of heterogeneity, not the binary reject outcome: per-model gap ranges span $5.4\text{--}13.9\times$ from the best- to worst-calibrated subject, and across the six larger models the *same* subjects sit at each extreme.

Subjects that are robustly worst-calibrated. Across the six larger models (excluding Qwen-0.5B due to its positional-bias confound), two subjects appear in the top-10 most miscalibrated for every model in our study: `virology` and `professional_law`. A consistent cluster of college-level technical subjects—`abstract_algebra`, `college_mathematics`, and `machine_learning`—each appears in the top-5 most miscalibrated for four of the six larger models. Together these form two recurring patterns: (i) college-level quantitative and technical content and (ii) niche professional or factual knowledge (`virology`, `professional_law`)—the high-stakes domains where systematic overconfidence creates the largest harm potential.

Subjects that are robustly best-calibrated. Symmetrically, `marketing` and `high_school_psychology` appear in the bottom-10 best-calibrated for every model in our study, and `high_school_government_and_politics` is consistently among the best-calibrated across models. Two patterns are visible. First, MMLU’s high-school-level subjects systematically show smaller calibration gaps than their college-level counterparts—the level of specialization required appears to track miscalibration. Second, the bottom-10 intersection is dominated by Social Sciences and the “Other” (applied/factual) MMLU category, with no STEM subject appearing.

Confidence vs. accuracy by entropy quartile per model (shaded = calibration gap)




Figure 3: Confidence (solid) and accuracy (dashed) by entropy quartile per model, with the shaded gap. The pattern progresses systematically across the Qwen family: at small scale (0.5B) the gap shrinks with rising entropy; at intermediate scale (1.5B) the curves fall in parallel; at larger scales (7B, 14B) the gap widens with entropy as confidence stays near saturation across the first three quartiles while accuracy still falls with entropy.

5.5 Sensitivity Analysis

Sensitivity-analysis conclusions extend uniformly: ECE rankings are stable across bin counts (Spearman ≥ 0.97 vs. 20-bin baseline), the Negative Log Calibration Score (a strictly proper scoring rule) ranks subjects nearly identically (Spearman ≥ 0.79), and two-level vs. three-level ICC estimates are numerically identical (consistent with $\sigma_{\text{domain}}^2 = 0$). The FDR rejection is also robust to threshold: at $\alpha = 0.01$, the four largest models reject 55–57/57 and the smaller Qwens 46–55/57 (the 4-bit 1.5B run is the most sensitive, discussed in the precision-effect appendix).

5.6 Cross-Model Agreement (RQ2)

Figure 5 shows the 7×7 Spearman rank correlation matrix for subject-level ECE across all seven main models. Among the six larger models, all 15 pairwise correlations fall in $[0.80, 0.94]$. The within-family pair Qwen-7B vs Qwen-14B is the highest at 0.94; cross-family GPT-4o-mini vs Qwen-7B is 0.92 and GPT-4o-mini vs Qwen-14B is 0.91.

Qwen-0.5B is the conspicuous exception, with correlations $[0.13, 0.20]$ to the other six models (bootstrap 95% CIs include zero). This is consistent with its $10.6\times$ positional bias: the smallest model’s miscalibration is dominated by a spurious preference for the D position rather than genuine knowledge gaps. We interpret this as evidence that, below approximately 1B parameters, positional bias dominates the calibration signal in our elicitation protocol.




Figure 4: Mean calibration gap per subject, sorted ascending, for each of the seven main models. Red bars indicate subjects where H_0 : mean gap = 0 is rejected at FDR $\alpha = 0.05$; gray bars are not rejected.

6 Discussion

Where calibration succeeds versus fails. The two strongest findings are descriptive but coherent. Calibration is reliably good on high-school-level Social Sciences and applied/factual subjects (`marketing`, `high_school_psychology`, `high_school_government_and_politics`), and reliably bad on college-level quantitative and technical content (`abstract_algebra`, `college_mathematics`, `machine_learning`) and on niche professional and factual knowledge (`virology`, `professional_law`). The same patterns appear in every model larger than 1B parameters that we examined: all 15 pairwise Spearman correlations among the six larger models fall in $[0.80, 0.94]$, with cross-family GPT-4o-mini vs Qwen-7B/14B at 0.92/0.91 comparable to the strongest within-family pairs. Subject-level miscalibration is therefore primarily a property of the task, providing empirical motivation for scoping corrections at the subject level rather than globally—though each correction would still be fit per-model on held-out data.

Within-family scale. Within the Qwen family at matched precision (1.5B fp16, 3B, 7B, 14B) the entropy coefficient on the calibration gap grows monotonically with scale and the between-subject ICC jumps from 1.5B to 3B then plateaus, while the mean-overconfidence intercept peaks at 3B before declining as accuracy catches up. The qualitative picture—capable models maintaining near-saturated confidence as items become harder—is consistent with prior reports of RLHF-induced confidence saturation [Kadavath et al., 2022, Nakkiran et al., 2025, Yaldiz et al., 2026]. We treat this as an observation about four checkpoints rather than a controlled mechanism test.

Implications for recalibration. The audit is the diagnostic arm of a broader recalibration picture; the operational arm is studied directly by Luo et al. [2025], who fits subject-specific temperature parameters per MMLU subject and shows improvements over a global correction. Our contribution is the upstream question—whether subject-level scoping is empirically motivated—which the cross-model stability above answers affirmatively. The robustly worst-calibrated subjects also serve as a transferable diagnostic for flagging miscalibration on new models before deployment.

Limitations. Domain-level variance is estimated at zero with only four domain groups, so we cannot cleanly partition between-subject variance into domain and subject-within-domain components. The mixed model assumes homoskedastic residuals, which is violated when calibration




Figure 5: Spearman rank correlation of subject-level ECE across the seven main models (57 common subjects). Among the six larger models, all 15 pairwise correlations fall in $[0.80, 0.94]$. Qwen-0.5B’s correlations with the others are $[0.13, 0.20]$, reflecting a qualitatively different miscalibration pattern driven by extreme positional bias.

gap variance depends on per-subject accuracy. Confidence is computed from logprobs at a single token position, which may not faithfully represent the model’s internal uncertainty after RLHF [Kadavath et al., 2022], and results are specific to MMLU’s four-choice format. The within-family scaling comparison is observational; the Qwen-0.5B endpoint is contaminated by extreme positional bias and does not contribute usable signal. 4-bit quantization affects five of the seven models; a precision comparison at 1.5B (4-bit vs. fp16) showed accuracy $0.536 \rightarrow 0.571$ and FDR rejection $46/57 \rightarrow 55/57$ at $\alpha = 0.01$, suggesting quantization amplifies positional priors and depresses confidence at the smallest scales (we use the fp16 1.5B as canonical and treat the rest as a precision-effect appendix); we did not run the comparison at larger scales.

Future directions. Additional Qwen checkpoints (e.g., 32B) would extend the scale progression beyond 14B; a third model family (Mistral or Gemma at matched scale) would strengthen the cross-family claim; and applying the framework to MMLU-Pro would test whether the high-school-vs-college gap persists at higher difficulty.

7 Conclusion

Subject-level calibration heterogeneity in LLMs is large and stable across model families above approximately 1B parameters. Applied/factual high-school subjects are reliably well-calibrated; college-level quantitative and technical content and niche professional and factual knowledge are

reliably bad. Subject-level ECE rankings correlate strongly across all 15 pairs of larger models (Spearman $\rho \in [0.80, 0.94]$), pointing to task structure rather than any particular training pipeline. Within the Qwen family the entropy coefficient on the calibration gap grows monotonically with scale, qualitatively in line with prior reports of RLHF-induced confidence saturation. The audit’s stable targets provide empirical motivation for subject-aware over global recalibration scoping [Luo et al., 2025] and supply a transferable diagnostic for flagging miscalibration on new models before deployment.

Statement of Work

Data analysis. Jennifer Zhang conducted the original two-model analysis (GPT-4o-mini, Llama-3-8B-Instruct), including the multilevel model, calibration curves, FDR correction, and sensitivity checks. John Wright collected and analyzed the four additional Qwen-2.5 model runs (0.5B, 1.5B 4-bit, 1.5B fp16, 3B, 7B, 14B), extended the multilevel and cross-model framework to seven models, and produced the within-family scale, entropy-mechanism, and quantization-effect analyses. **Coding.** Jennifer Zhang built the original pipeline; John Wright extended it to support arbitrary models and produced the parameterized inference notebook used for the open-weight models. **Writing.** Jennifer Zhang drafted the original report; John Wright revised it for the seven-model expansion and produced the final manuscript. **Presentation.** Jennifer Zhang prepared the slide deck.

References

- Solon Barocas, Moritz Hardt, and Arvind Narayanan. *Fairness and Machine Learning: Limitations and Opportunities*. MIT Press, 2023.
- Yoav Benjamini and Yosef Hochberg. Controlling the false discovery rate: A practical and powerful approach to multiple testing. *Journal of the Royal Statistical Society, Series B*, 57(1):289–300, 1995.
- Tim Dettmers, Mike Lewis, Younes Belkada, and Luke Zettlemoyer. LLM.int8(): 8-bit matrix multiplication for transformers at scale. In *Advances in Neural Information Processing Systems*, 2022.
- Bradley Efron. *Large-Scale Inference: Empirical Bayes Methods for Estimation, Testing, and Prediction*. Cambridge University Press, 2012.
- Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. On calibration of modern neural networks. In *Proceedings of the 34th International Conference on Machine Learning*, 2017.
- Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. Measuring massive multitask language understanding. In *Proceedings of the International Conference on Learning Representations*, 2021.
- Zhengbao Jiang, Jun Araki, Haibo Ding, and Graham Neubig. How can we know when language models know? On the calibration of language models for question answering. *Transactions of the Association for Computational Linguistics*, 9:962–977, 2021.
- Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas Schiefer, Zac Hatfield-Dodds, Nova DasSarma, Eli Tran-Johnson, et al. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*, 2022.

- Beier Luo, Shuoyuan Wang, Yixuan Li, and Hongxin Wei. Your pre-trained LLM is secretly an unsupervised confidence calibrator. *arXiv preprint arXiv:2505.16690*, 2025.
- Preetum Nakkiran, Arwen Bradley, Adam Goliński, Eugene Ndiaye, Michael Kirchhof, and Sinead Williamson. Trained on tokens, calibrated on concepts: The emergence of semantic calibration in LLMs. *arXiv preprint arXiv:2511.04869*, 2025.
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. In *Advances in Neural Information Processing Systems*, volume 35, pages 27730–27744, 2022.
- Stephen W. Raudenbush and Anthony S. Bryk. *Hierarchical Linear Models: Applications and Data Analysis Methods*. Sage Publications, 2nd edition, 2002.
- Hexiang Tan, Wanli Yang, Junwei Zhang, Xin Chen, Rui Tang, Du Su, Jingang Wang, Yuanzhuo Wang, Fei Sun, and Xueqi Cheng. BaseCal: Unsupervised confidence calibration via base model signals. *arXiv preprint arXiv:2601.03042*, 2026.
- Miao Xiong, Zhiyuan Hu, Xinyang Lu, Yifei Li, Jie Fu, Junxian He, and Bryan Hooi. Can LLMs express their uncertainty? An empirical evaluation of confidence elicitation in LLMs. In *Proceedings of the Twelfth International Conference on Learning Representations*, 2024.
- Duygu Nur Yaldiz, Evangelia Spiliopoulou, Zheng Qi, Siddharth Varia, Srikanth Doss, and Nikolaos Pappas. Balancing classification and calibration performance in decision-making LLMs via calibration aware reinforcement learning. *arXiv preprint arXiv:2601.13284*, 2026.
- Bianca Zadrozny and Charles Elkan. Transforming classifier scores into accurate multiclass probability estimates. In *Proceedings of the Eighth ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 694–699, 2002.